

Data Ingestion from the RDS to HDFS using Sqoop

Sqoop Import command used for importing table from RDS to HDFS:

The command to import data from RDS to HDFS is:

```
1. sqoop import --connect <RDS connection string> \  
2. --username <RDS Username> --password <RDS Password> \  
3. --table <Table name in RDS> \  
4. --null-string '\\N' --null-non-string '\\N' \  
5. --target-dir <path in HDFS where the data is imported to> \  
6. -m <no. of mapper jobs> --as-parquetfile
```

Here,

RDS connect string: "jdbc:mysql://upgradtest.cyaielc9bmnf.us-east-1.rds.amazonaws.com/testdatabase"

Username: "student"

Password: "STUDENT123"

Table name: "SRC_ATM_TRANS" and

Target dir: "/user/root/atm_trans"

We are saving this as a "parquet" file and running only 1 mapper job. So the sqoop import command will be as given below:

```
sqoop import \  
--connect jdbc:mysql://upgradtest.cyaielc9bmnf.us-east-1.rds.amazonaws.com/testdatabase \  
--table SRC_ATM_TRANS \  
--username student --password STUDENT123 \  
--null-string '\\N' --null-non-string '\\N' \  
--target-dir /user/root/atm_trans \  
-m 1 --as-parquetfile
```

```

File System Counters
  FILE: Number of bytes read=0
  FILE: Number of bytes written=181302
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=41970
  HDFS: Number of bytes written=44677613
  HDFS: Number of read operations=68
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=10
Job Counters
  Launched map tasks=1
  Other local map tasks=1
  Total time spent by all maps in occupied slots (ms)=65219
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=65219
  Total vcore-milliseconds taken by all map tasks=65219
  Total megabyte-milliseconds taken by all map tasks=66784256
Map-Reduce Framework
  Map input records=2468572
  Map output records=2468572
  Input split bytes=87
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=1063
  CPU time spent (ms)=67100
  Physical memory (bytes) snapshot=637370368
  Virtual memory (bytes) snapshot=2817974272
  Total committed heap usage (bytes)=478674944
File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=0
21/04/01 16:29:32 INFO mapreduce.ImportJobBase: Transferred 42.6079 MB in 87.262
7 seconds (499.9903 KB/sec)
21/04/01 16:29:32 INFO mapreduce.ImportJobBase: Retrieved 2468572 records.

```

We see that all the records (2468572) are successfully imported to HDFS.

Command used to see the list of imported data in HDFS:

Below is the command to see the files present in the directory atm_trans, i.e., the HDFS location to which we imported the data:

```
hadoop fs -ls /user/root/atm_trans
```

Where “/user/root/atm_trans” is the location where the data is imported in to the HDFS and “99b03bca-c9e7-4fb6-916f-1baa29e96797.parquet” is the record. So, we pass the full path in the “hadoop fs -ls” command. The data imported looks as shown below.

```
[root@ip-10-0-0-208 ~]# hadoop fs -ls /user/root/atm_trans
Found 3 items
drwxr-xr-x  - root supergroup          0 2021-04-01 16:28 /user/root/atm_trans/
.metadata
drwxr-xr-x  - root supergroup          0 2021-04-01 16:29 /user/root/atm_trans/
.signals
-rw-r--r--  3 root supergroup  44667278 2021-04-01 16:29 /user/root/atm_trans/
99b03bca-c9e7-4fb6-916f-1baa29e96797.parquet
```

Screenshot of the imported data:

Let us check the parquet file that was created during sqoop import and see if the required file is present in it.

```
[root@ip-10-0-0-208 ~]# hadoop fs -ls /user/root/atm_trans/99b03bca-c9e7-4fb6-916f-1baa29e96797.parquet
-rw-r--r--  3 root supergroup  44667278 2021-04-01 16:29 /user/root/atm_trans/
99b03bca-c9e7-4fb6-916f-1baa29e96797.parquet
```